LAB 2 discussion

Simulation based Process Scheduling

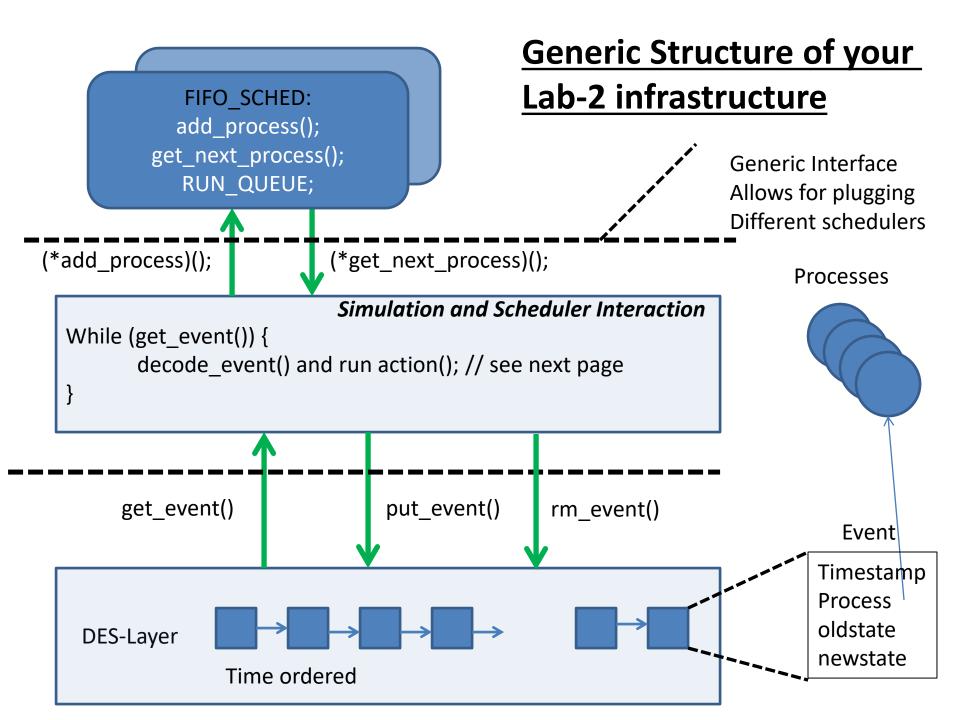
- Explore multiple Schedulers and what OS tracks
- You will learn how to write a Discrete Event Simulator
- Specification of Process and Random Behavior

```
Process-1
                                200
                                             40
                                                         90
                     40
                                100
                                             10
                                                         40
                                  2.0
                     50
                                              10
                                                         10
Process-4
                                200
                                                         20
                     60
                                                         10
                               Total
                                              CPU
              process
              arrival/start
                               CPU-time
                                                         Burst
                                             Burst
```

- Cpu-burst = random (range [1.. Proc.cpu_burst])
- Ioburst = random (range [1.. Proc.io_burst])

What you will generate?

```
cat input4
   200
         40
              90
                             static-prio Finish Time Turnaround time
   100
              40
40
         10
50
    20
        10
              10
60
   200
              20
$ ./sched -sF input4 rfile
FCFS
0000:
          0
              200
                     40
                           90
                                      575
                                             575
                                                    360
                                                            92
0001:
              100
                           40
                                      532
                                             492
                                                    300
         40
                     10
0002:
                                              32
         50
               20
                     10
                           10
                                       82
                                                     10
0003:
         60
              200
                      5
                           20
                                     1174
                                            1114
                                                    773
                                                           141
           44.29 81.26
                         553.25
                                  62.50 0.341
```



```
void Simulation() {
         EVENT* evt;
         while( (evt = get event()) ) {
                 Process *proc = evt->evtProcess; // this is the process the event works on
                 CURRENT TIME = evt->evtTimeStamp;
                 int transition = evt->transition; // for accouting
                 int timeInPrevState = CURRENT TIME - proc->state ts;
                 delete evt; evt = nullptr; // remove cur event obj and don't touch anymore
                 switch (transition) { // encodes where we come from and where we go
                 case TRANS TO READY:
                         // must come from BLOCKED or CREATED
                         // must add to run queue, no event created
                         CALL SCHEDULER = true;
                         break;
                 case TRANS TO PREEMPT: // similar to TRANS TO READY
                         // must come from RUNNING (preemption)
                         // must add to run queue, no event created
                         CALL SCHEDULER = true;
                         break;
                 case TRANS TO RUN:
                         // create event for either preemption or blocking
                         break;
                 case TRANS TO BLOCK:
                         //create an event for when process becomes READY again
                         CALL SCHEDULER = true;
                         break;
                 if (CALL SCHEDULER) {
                         if (get next event time() == CURRENT TIME)
                                 continue; //process next event from Event queue
                         CALL SCHEDULER = false; // reset global flag
                         if (CURRENT RUNNING PROCESS == nullptr) {
                                 CURRENT RUNNING PROCESS = THE SCHEDULER->get next process();
                                 if (CURRENT RUNNING PROCESS == nullptr)
                                      continue;
                                 // create event to make this process runnable for same time.
```

./sched -v -e input_show rfile

Input file

0 100 10 10 20 100 20 10

ShowEventQ: 0:0 20:1

0 0 0: CREATED -> READY

AddEvent(0:0:RUNNG): 20:1:READY ==> 0:0:RUNNG 20:1:READY

0 0 0: READY -> RUNNG cb=8 rem=100 prio=1

AddEvent(8:0:BLOCK): 20:1:READY ==> 8:0:BLOCK 20:1:READY

8 0 8: RUNNG -> BLOCK ib=2 rem=92

AddEvent(10:0:READY): 20:1:READY ==> 10:0:READY 20:1:READY

10 0 2: BLOCK -> READY

AddEvent(10:0:RUNNG): 20:1:READY ==> 10:0:RUNNG 20:1:READY

10 0 0: READY -> RUNNG cb=10 rem=92 prio=1

AddEvent(20:0:BLOCK): 20:1:READY ==> 20:1:READY 20:0:BLOCK

20 1 0: CREATED -> READY

20 0 10: RUNNG -> BLOCK ib=7 rem=82

AddEvent(27:0:READY): ==> 27:0:READY

AddEvent(20:1:RUNNG): 27:0:READY ==> 20:1:RUNNG 27:0:READY

20 1 0: READY -> RUNNG cb=7 rem=100 prio=3

AddEvent(27:1:BLOCK): 27:0:READY ==> 27:0:READY 27:1:BLOCK

RESULTS OF SIMULATION

FCFS

0000: 0 100 10 10 2 | 234 234 89 45

0001: 20 100 20 104 | 226 206 77 29

SUM: 234 85.47 57.26 220.00 37.00 0.855

./sched -v -e input show rfile

ShowEventQ: 0:0 20:1

0 0 0: CREATED -> READY

line triggered by "-v"
Timestamp pid howlong:
FROM -> TO

Input file

0 100 10 10 20 100 20 10

0 0 0: READY -> RUNNG cb=8 rem=100 prio=1

AddEvent(8:0:BLOCK): 20:1:READY ==> 8:0:BLOCK 20:1:READY

AddEvent(0:0:RUNNG): 20:1:READY ==> 0:0:RUNNG 20:1:READY

One Event

8 0 8: RUNNG -> BLOCK ib=2 rem=92

AddEvent(10:0:READY): 20:1:READY ==> 10:0:READY 20:1:READY

10 0 2: BLOCK -> READY

AddEvent(10:0:RUNNG): 20:1:READY ==> 10:0:RUNNG 20:1:READY

10 0 0: READY -> RUNNG cb=10 rem=92 prio=1

AddEvent(20:0:BLOCK): 20:1:READY ==> 20:1:READY 20:0:BLOCK

20 1 0: CREATED -> READY

20 0 10: RUNNG -> BLOCK ib=7 rem=82

AddEvent(27:0:READY): ==> 27:0:READY

AddEvent(20:1:RUNNG): 27:0:READY ==> 20:1:RUNNG 27:0:READY

20 1 0: READY -> RUNNG cb=7 rem=100 prio=3

AddEvent(27:1:BLOCK): 27:0:READY ==> 27:0:READY 27:1:BLOCK

line triggered by "-e"
event added: (time,pid,transition)
EventQ-Before → EventQ-After
(should be time ordered)

./sched -v -t input_show rfile

Input file

0 100 10 10 0.00: CREATED -> READY SCHED (1): 0:0 20 100 20 10 0 0 0: READY -> RUNNG cb=8 rem=100 prio=1 8 0 8: RUNNG -> BLOCK ib=2 rem=92 SCHED (0): 10 0 2: BLOCK -> READY SCHED (1): 0:10 10 0 0: READY -> RUNNG cb=10 rem=92 prio=1 20 1 0: CREATED -> READY 20 0 10: RUNNG -> BLOCK ib=7 rem=82 SCHED (1): 1:20 line triggered by "-t" to show scheduler runqueue 20 1 0: READY -> RUNNG cb=7 rem=100 prio=3 27 0 7: BLOCK -> READY SCHED(len): { "pid:timestamp" }* 27 1 7: RUNNG -> BLOCK ib=9 rem=93 SCHED (1): 0:27 27 0 0: READY -> RUNNG cb=7 rem=82 prio=1 where len is length of RQ followed by RQ entries 34 0 7: RUNNG -> BLOCK ib=1 rem=75 SCHED (0): 35 0 1: BLOCK -> READY SCHED (1): 0:35 Current cpu-burst 35 0 0: READY -> RUNNG cb= rem=75 prio=1 36 1 9: BLOCK -> READY Current io-burst 44 0 9: RUNNG -> BLOCK ib=9 em=ob SCHED (1): 1:36 Dynamic priority 44 1 8: READY -> RUNNG cb=11 rem=93 prio=3 53 0 9: BLOCK -> READY Remaining cputime 55 1 11: RUNNG -> BLOCK ib=4 rem=82

./sched –p input_show rfile

To help with the preemption (EPRIO)

There are 2 conditions that need to be met:

Cond1: unblocked_process.dynprio > current.dynprio AND

Cond2: next event for current proc is > current time

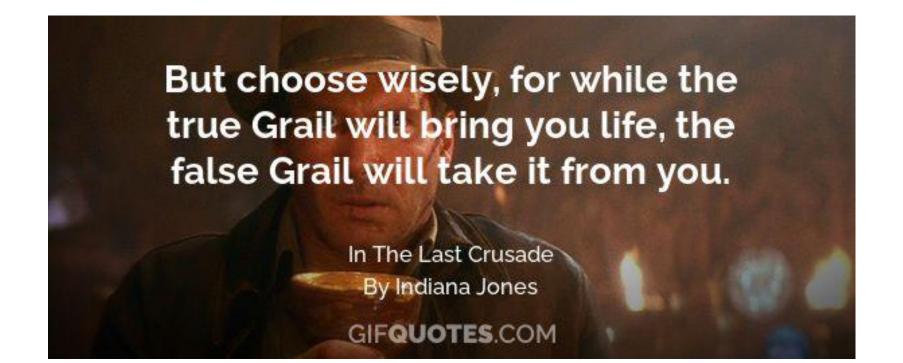
Time till next event of current running process

Accounting

- Make sure you do backward accounting (how long was process in a particular state.
- Forward accounting (how long will I be in a state) will not work.
- Charge the right bucket (cpu, blocked, ready)
 when you get out of its state, not when you get in (otherwise "E"
 scheduler won't work and forward accounting in real systems is not
 possible)
- You have all the information how long you were in the state you are transitioning out of.

Data Structures

- Schedulers are often about pooling requests and selecting:
 - Queueing, searching, indexing, mapping, priorities
 - Efficiency matters, not just effectiveness
 - Underlying data structures matter



Example: Queue Semantics

- · Assume existing "queue" size N
- Consider the following operations
 - Add front / delete tail (FT)
 - Add Tail / delete front (TF)
 - Iterate through queue with
 - C++ iterator (Iter) or C++ [] operator
- Time these operations
- Consider the following container classes:
 - Deque (Q)
 - Vector (V)
 - List (L)

Extract behavior

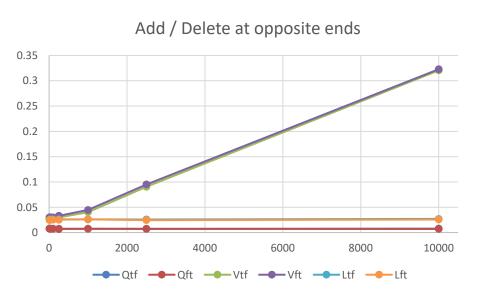
• https://dev.to/pratikparvati/c-stlcontainers-choose-your-containerswisely-4|c4 excellent article with good graphics

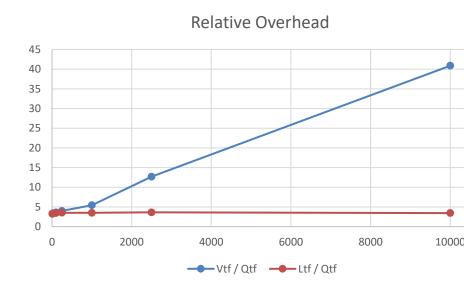
 https://www.geeksforgeeks.org/containe rs-cpp-stl/

```
eventQ = list<Event*> because we add in the middle and remove at the front procqueue = deque<Process*> because we push and pop at the front and end srtfqueue = list<Process*> because we add in the middle and remove at the front prioqueue = vector<deque<Process*>>
```

Add / Delete Opposite ends

- Scale the size of the queue (x-axis = N)
- TF and FT semantics
- Why are vectors so expensive?

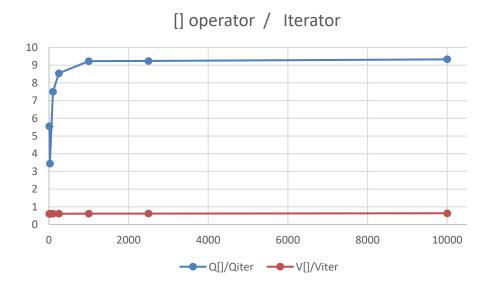




Iterations

• Iterate by 2 approaches:

```
iter: for ( x : objs ) { x } ;
[]: for (i = 0; i<obj.size(); i++) { obj[i]; }</pre>
```

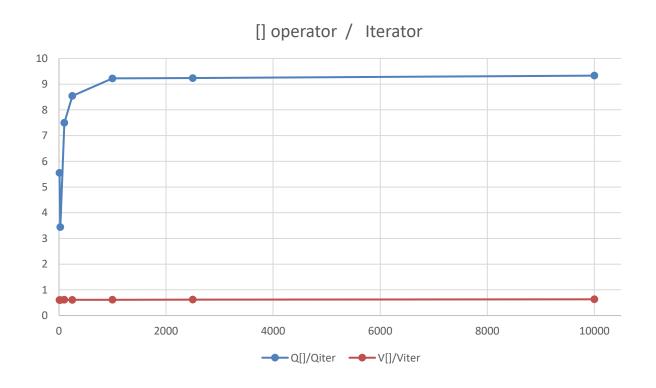


note: list[] does not

exist

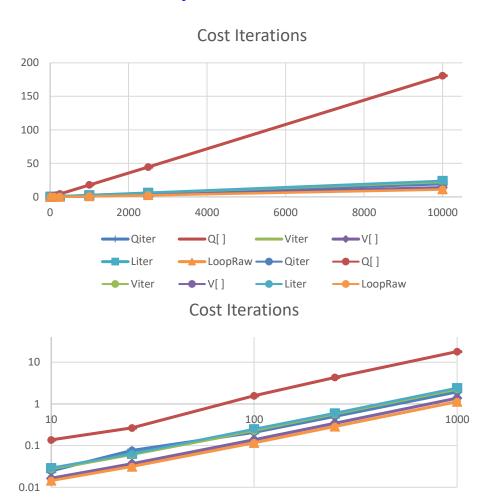
Iteration vector vs deque

• [] operator relatively expensive for deque over iterator, but not for vector



Iteration vector vs deque vs list

- Avoid [] operator on queues
- Think about underlying data structures how something would be implemented and think about cost of operations.



---- Qiter

--Q[]

→V[]

— Q[] **—** LoopRaw

Moral of the story

- Don't just use STL data structures
- Think how they are being used
- Think a bit about implementation and efficiency